

Comparing Safety-Removal Subspaces in Aligned LLMs

Evan Scamehorn

University of Wisconsin
scamehorn@wisc.edu

Adam Venton

University of Wisconsin
venton@wisc.edu

Calvin Kosmatka

University of Wisconsin
ckosmatka@wisc.edu

Kyle Sha

University of Wisconsin
kasha2@wisc.edu

Zeke Mackay

University of Wisconsin
afmackay@wisc.edu

Abstract

Modern Large Language Models (LLMs) undergo extensive alignment to ensure they don't produce harmful outputs, while still being as helpful as possible. However, these aligned models are fragile and can be jailbroken by a variety of methods. Recent research suggests that this fragility may stem from compact safety mechanisms whose relationship to general utility remains poorly understood. Several methods have been proposed to isolate or remove safety-related structure, including Difference in Means (DIM), Refusal Cone Optimization (RCO), and ActSVD. Conversely, other research suggests that it might be impossible to completely linearly separate safety and utility in current LLMs. In this project, we compare DIM and ActSVD on LLaMA-3.1-Instruct using harmful and harmless instruction datasets, and we add a Geometry-paper-style Representational Independence (RepInd) profile analysis for DIM-derived directions. We evaluate behavioral safety and utility with JailbreakBench attack success rate, harmless instruction compliance, and perplexity. We then evaluate geometric agreement between DIM and ActSVD using Mode Subspace Overlap (MSO) and directly compare DIM safety directions against harmless-instruction utility PCA subspaces. This framing answers a concrete question: do current linear methods recover the same refusal mechanism, how much do safety directions overlap with utility activations, and do nearby refusal directions remain causally independent under ablation?

1 Introduction

1.1 Context and Motivation

As LLMs have become more powerful and more accessible, LLM alignment techniques have become significantly more important. However, despite much research in this area, modern LLMs remain fragile, being able to be jailbroken by a variety of methods, including special prompts and

white-box methods like DIM (Arditi et al., 2024) and ActSVD (Wei et al., 2024). Improving the robustness of these models will require moving beyond current empirical methods and developing a deep theoretical understanding of how safety and utility are represented within models.

1.2 The Problem: Superposition in Activation Space

Recent research into mechanistic interpretability has shown that model behavior is determined by distinct directions within the activation space. For example, there might be a single direction that determines model refusal, and by adjusting that direction within the activation space, we can control whether or not the model refuses to answer a prompt (Arditi et al., 2024). However, a fundamental challenge to this way of understanding activation spaces is *superposition*. When models need to represent more features than they have dimensions, some dimensions must contain information about multiple features. Recent research suggests that safety and utility share representation capacity, and thus any attempt to adjust one of these features through linear modification may (and probably will) degrade the other.

1.3 The Gap in Current Literature

Recent research has introduced several methods to identify safety and utility subspaces. Difference in Means (DIM) (Arditi et al., 2024) identifies a single vector mediating refusal, Refusal Cone Optimization (RCO) maps a multidimensional cone space (Wollschläger et al., 2025), and ActSVD isolates low-rank matrices via singular value decomposition (SVD) (Wei et al., 2024). While each of these methods successfully creates a safety-related subspace, each has a different mathematical geometry, and it is not well understood how they all relate.

1.4 Proposed Research

In this project, we investigate safety representations through three connected questions. First, do DIM and ActSVD identify the same safety-removal mechanism, or do they modify distinct regions of the model? Second, how much does the DIM safety direction overlap with harmless-instruction utility activations? Third, when these methods weaken refusal behavior, how much general utility do they preserve? We implement DIM and ActSVD on LLaMA-3.1-Instruct and add a local RepInd profile analysis based on the Geometry paper. Full optimized RCO cone training remains an extension path. We evaluate behavioral effects with JailbreakBench attack success rate, harmless Alpaca compliance, and Pile/Alpaca perplexity. We then compute overlap between extracted subspaces using Mode Subspace Overlap (MSO). We also evaluate Representational Independence for DIM-derived directions by measuring how each direction’s layerwise cosine profile changes when another direction is ablated.

Comparing DIM and ActSVD tests cross-method agreement between safety-removal mechanisms. To answer safety-utility overlap directly, we also compute a matched activation-space analysis: DIM safety directions are compared against the top PCA directions of harmless instruction activations at the same layers and token positions. This gives a direct MSO estimate of how much safety lies inside the utility activation subspace.

2 Literature Survey

Arditi et al. (2024) show that for several common chat models, a manipulation of a single dimensional subspace is enough to both induce refusal of non-harmful requests, and turn off the refusal of harmful requests. They use a procedure called **difference-of-means** to identify this subspace. This technique measures the average activations at each token position of each transformer layer, compared between a set of harmful requests and a set of harmless requests. Once these difference-of-means vectors are constructed, they score each for its ability to cause the model to refuse harmless requests and respond to harmful requests. The vector with the highest score is normed and selected as the refusal dimension vector, denoted by $\hat{r}$. This refusal dimension can then be ablated in the residual stream at inference time to bypass refusal behavior. This vector can also be used to create

a jailbroken model by updating model weights according to the following formula: $W'_{\text{out}} \leftarrow W_{\text{out}} - \hat{r}\hat{r}^T W_{\text{out}}$. They find that this attack method is successful on the Qwen model family with and without the system prompt, and on the Llama model family without the system prompt, and has little effect on model coherence.

Wei et al. (2024) introduce low-rank decomposition methods designed to identify specific ranks within a weight matrix related to given LLM behaviors. Their ActSVD algorithm performs Singular Value Decomposition on the product of the model weights and input activations ($W X_{\text{in}}$), and yields an orthogonal projection matrix (Π).

Removing the top safety-critical ranks ActSVD identifies causes the model to completely stop rejecting unsafe prompts, and the model’s utility is severely compromised. These findings suggest that safety regions in aligned models are also crucial for its general utility. To disentangle safety from utility, the authors remove safety ranks orthogonal to utility ranks using $\Delta W = (I - \Pi^u)\Pi^s W$. This yields higher attack success rate for unsafe prompts while maintaining zero-shot accuracy for utility prompts.

The fact that naively ablating safety ranks destroys utility, whereas surgically removing disentangled ranks preserves it, indicates that top safety ranks and top utility ranks heavily overlap. The necessity of this orthogonal projection matrix provides strong evidence against the hypothesis of strict linear separability between safety and utility. Ultimately, ActSVD provides rank-level evidence for superposition: safety and utility share representational capacity and are not linearly distinct.

Ponkshe et al. (2026) demonstrates fundamental mathematical limitations of linear subspace-based safety defenses, arguing that safety is not linearly separable from utility. Removing any specific safety-related subspace inherently degrades the overall performance of the model. The authors support this hypothesis through a series of empirical evaluations. Using singular value decomposition and mode subspace overlap, the study reveals that the principal directions amplifying safe behaviors also amplify useful ones, indicating that these directions do not constitute a distinctly separable safety subspace. Furthermore, attempts to mitigate harmful behaviors via orthogonal projection resulted in a proportional drop in the model’s utility. The researchers also found

that providing the model with helpful and harmful inputs produced highly overlapping activations. Collectively, these findings further challenge the hypothesis that safety and utility operate within linearly separable subspaces.

Wollschläger et al. (2025) generalize the identification of safety subspaces to conic regions of multiple basis refusal vectors as opposed to one refusal direction. Instead of testing pairs of harmful and harmless prompts, their methods of Refusal Direction Optimization and Refusal Cone Optimization perform gradient descent to converge on refusal vector direction(s). They leverage two properties of ideal refusal vectors in loss functions for optimization:

- Given a refusal direction r , scalar α , initial activation x_i , and revised activation $\tilde{x}_i = x_i + \alpha \cdot r$, refusal probability should scale with α .
- Removing the refusal direction should not affect harmless prompts while allowing response to harmful prompts.

This research finds significant jailbreaking performance gains using one refusal direction, with further gains up to a four-dimensional refusal region. Testing on Gemma-2, Llama 3, and Qwen 2.5 model families and benchmarks such as TruthfulQA, ablation of multiple refusal vectors is shown to have better attack success and lower side-effects on model performance.

Furthermore, noting that vector orthogonality does not guarantee causal independence, the research defines a notion of representational independence between vectors, in which the ablation of one direction does not impact the effects of the other. Ablating three or more representationally independent refusal vectors is found to have higher attack success than difference-of-means direction ablation. This further shows that safety and utility occupy a complex subspace within LLMs.

3 Methodology

Targeting Llama-3.1-8B-Instruct, we compare three methods for extracting safety subspaces:

Method	Operates on	Intervention
DIM	Residual stream activations	Inference-time ablation by 1 direction
ActSVD	Weight matrices and activations per linear layer	Removal of safety ranks from weight matrix
RCO	Residual stream activations	Inference-time ablation by subspace of 2 basis vectors

Table 1: The three methods operate at different levels of the model but target the same behavioral outcome (removing refusal).

3.1 Safety Subspace Identification

We implement three complementary methods for extracting safety-relevant subspaces from the model’s internal representations.

Difference-in-Means (DIM). Following Arditi et al. (2024), we compute the mean residual stream activations at each layer l and post-instruction token position i for sets of harmful and harmless prompts. The difference-in-means vector is $\mathbf{r}_i^{(l)} = \boldsymbol{\mu}_i^{(l)} - \mathbf{v}_i^{(l)}$, where $\boldsymbol{\mu}$ and $\mathbf{v}$ are the mean activations over harmful and harmless prompts, respectively. We select the single most effective vector $\hat{\mathbf{r}}$ by evaluating each candidate’s ability to bypass refusal when ablated and to induce refusal when added. The selected unit-norm vector defines a one-dimensional safety subspace.

ActSVD Safety and Utility Ranks. Following Wei et al. (2024), we perform Singular Value Decomposition on the product of model weights and input activations $W\mathbf{X}_{\text{in}}$ for both safety and utility calibration datasets, yielding $USV^\top \approx W\mathbf{X}_{\text{in}}$. The orthogonal projection matrices $\Pi^s = U^s(U^s)^\top$ and $\Pi^u = U^u(U^u)^\top$ project onto the top r^s safety and top r^u utility rank subspaces, respectively. To disentangle safety from utility, we compute the isolated safety projection: $\Delta W(r^u, r^s) = (I - \Pi^u)\Pi^s W$. While Wei et al. evaluate on Llama-2 7B/13B, the method operates on generic linear layers and transfers directly to Llama-3.1 8B.

Refusal Cone Optimization (RCO). Following Wollschläger et al. (2025), the Geometry codebase can use gradient-based optimization to discover multiple refusal directions that together form a multi-dimensional conic region. The optimization minimizes a composite loss encoding

two properties: (1) monotonic scaling of refusal probability with the magnitude of activation addition, and (2) surgical ablation that bypasses refusal on harmful prompts while preserving behavior on harmless prompts. A retain loss based on KL divergence ensures minimal side effects on harmless inputs. In our current results, full optimized RCO training is treated as an extension path; we do, however, run the RepInd profile test on DIM-derived directions.

3.2 Subspace Comparison

Our comparison phase addresses two questions. First, *cross-method consistency*: do DIM, ActSVD, and RCO converge on similar safety-relevant features, or does each capture a distinct aspect of the safety mechanism? Second, *safety-utility overlap*: how much does each layer’s safety direction lie inside a utility activation subspace? Third, *behavioral safety-utility tradeoff*: after removing safety behavior, how much does each method degrade useful behavior?

DIM and RCO operate in activation space, while ActSVD operates in weight space. We bridge these two regimes in two complementary ways. **Weight-delta MSO**: for each layer and weight type, we compute the thin SVD of $\Delta W = W_{\text{actsvd}} - W_{\text{base}}$ and project the DIM and RCO directions onto its left singular basis. This is a **capacity** measure: the realized perturbation $\Delta W x$ depends on the input distribution. **Activation-delta direction**: per-layer mean activation deltas computed across 64 harmless prompts provide a direct activation-space comparison between the base and ActSVD-modified models.

Mode Subspace Overlap (MSO). Following Punkshe et al. (2026), MSO measures the geometric overlap between two subspaces. For two matrices V and W , we extract their principal directions via thin SVD and select the smallest number of left singular vectors capturing an η -fraction of the energy. The MSO metric is defined as:

$$\text{MSO}(V, W; \eta) = \frac{\|S\|_F^2}{\min(k_V, k_W)}$$

where $S = Q_V^\top Q_W$ is the overlap matrix between the orthonormal bases. MSO ranges from 0 (orthogonal subspaces) to 1 (identical spans). We compute MSO between the DIM refusal direction and ActSVD weight-delta subspaces to assess cross-method agreement. We also compute MSO between DIM safety directions and matched utility PCA subspaces to quantify safety-utility entan-

glement directly. Because DIM yields a single direction while ActSVD yields multi-dimensional subspaces, cross-method MSO is bounded by the dimensionality asymmetry; we report the random baseline $\mathbb{E}[\text{overlap}] = \max(k_V, k_W)/d$ alongside each MSO value for calibration.

Representational Independence (RepInd).

Following Wollschläger et al. (2025), RepInd tests whether two individual directions are *causally* related, not merely geometrically similar. Two directions $\lambda, \mu \in \mathbb{R}^d$ are representationally independent if ablating one does not change the cosine similarity profile of the other across layers:

$$\forall l \in L : \cos(x^{(l)}, \lambda) = \cos(\tilde{x}_{\text{abl}(\mu)}^{(l)}, \lambda)$$

and vice versa. MSO may report high geometric overlap between directions that turn out to be causally independent, or low overlap between directions that are causally entangled via non-linear interactions across layers. In this codebase, the cones-repind implementation can train RDO directions, orthogonal directions, independent directions, and multi-vector refusal cones. Our project-owned RepInd script measures a direction’s layerwise cosine-similarity profile before and after ablating another direction. Without completed optimized RCO artifacts, we derive a small cone-like basis from high-norm DIM mean-difference candidates; with trained RCO artifacts, the same script can compare DIM, RDO, RepInd, and cone basis vectors directly.

4 Current Implementation Status

The current codebase has working DIM and ActSVD runs for LLaMA-3.1-Instruct, plus a direct safety-utility overlap analysis and a RepInd profile analysis. The DIM pipeline extracts a refusal direction, evaluates baseline/ablation/activation-addition completions, computes utility losses, and saves a modified model. The ActSVD pipeline computes low-rank safety and utility projections and saves a modified model after removing safety ranks orthogonal to utility ranks. The safety-utility overlap script computes per-layer DIM safety directions and compares them to top PCA bases of harmless instruction activations, producing per-layer, rank-sweep, and heatmap figures. The RepInd script computes cosine-profile changes before and after ablation for DIM and two DIM-derived cone-basis candidates. Full RCO training remains an extension because it requires a longer optimization run.

Our preliminary behavioral benchmark compares the base model, the DIM-modified model, and the ActSVD-modified model. DIM ablation increases JailbreakBench attack success rate from 0.16 to 1.00 while leaving harmless compliance at 1.00 and producing only small perplexity changes. ActSVD increases attack success rate to 0.63 but causes larger perplexity degradation. This suggests that DIM is a stronger safety-removal attack in our current setup, while ActSVD creates more collateral utility damage.



Figure 1: Preliminary safety-utility tradeoff. DIM removes refusal with little perplexity cost, while ActSVD partially removes refusal with higher utility degradation.

The first geometric comparison measures DIM-vs-ActSVD method overlap. MSO is close to the random baseline for most layers, with a mild local hotspot around layer 10. The second geometric comparison measures safety-vs-utility overlap directly by projecting DIM safety directions into harmless-instruction utility PCA subspaces. Together, these distinguish method agreement from the stronger safety-utility superposition question.




Figure 2: Average per-layer MSO between the DIM refusal direction and ActSVD weight-delta subspaces. Most layers are near the random baseline.

The direct safety-utility activation analysis shows substantially above-random overlap. With a rank-8 utility PCA subspace, mean safety-utility MSO is 0.192, compared to a random baseline of 0.00195. The selected DIM refusal direction at layer 11 has rank-8 utility-subspace overlap of 0.078. These values indicate that the safety direction is not geometrically isolated from the activation mani-

fold used by benign instructions, although overlap varies substantially by layer.




Figure 3: Per-layer overlap between DIM safety directions and rank-8 harmless-instruction utility PCA subspaces.




Figure 4: Mean safety-utility MSO increases as the utility PCA subspace rank grows, remaining above the random-subspace baseline.

The RepInd profile analysis adds the causal comparison from the Geometry paper. Using 32 harmful prompts, we compare DIM with two DIM-derived cone-basis candidates. Ablating DIM changes the first derived basis profile substantially (mean absolute profile change 0.227), while ablating that derived basis barely changes the DIM profile (0.0017). The second derived basis is nearly independent in both directions. This suggests that DIM occupies a dominant refusal component, while some additional high-norm candidate directions are much less causally coupled to DIM. Because these are DIM-derived candidates rather than fully optimized RCO vectors, this result should be read as a RepInd methodology run, not as a completed cone-optimization claim.




Figure 5: RepInd profile-change matrix for DIM and two DIM-derived cone-basis candidates. Rows are ablated directions; columns are measured directions. Lower off-diagonal values indicate greater representational independence.

5 Data Sets

We plan to use two primary datasets to conduct testing. Alpaca (Taori et al., 2023) will be used to test utility (refusal rate) on harmless fine-tuning data, while BeaverTails (Ji et al., 2023) will be used to test safety (attack success rate) against harmful fine-tuning data.

Alpaca is a dataset of 50,000+ common instructions and LLM-generated outputs. It is used most commonly in LLM fine-tuning. All tasks are harmless and should be completed without issue by an aligned Llama-3.1 model. The dataset contains three text fields:

- **instruction:** Description of the task assigned to the model.
- **input:** Additional information required to complete the task (e.g. initial code if the task is to refactor code).
- **output:** Answer to the task, generated by text-davinci-003 by OpenAI.

BeaverTails is a fine-tuning dataset with a focus on safety alignment. It contains question-answer pairs of a variety of harmful and harmless categories. Aligned Llama-3.1 models should be able to reasonably respond to safe prompts and decline to answer unsafe prompts. The dataset contains the following fields:

- **prompt:** Question provided to the LLM.
- **response:** Sample answer to the question from an aligned LLM.
- **category:** Harmful topic(s) ascribed to the question (e.g. animal abuse, hate speech, financial crime).
- **is_safe:** Whether the question provided is safe or unsafe (approximately 57% of questions are safe).

6 Data Analysis

Seeing as both testing datasets are purely textual and tests will be on compliance with prompts in datasets, little analysis can be done of either Alpaca or BeaverTails prior to their implementation in our models.

6.1 Alpaca

Alpaca contains instructions and LLM-generated outputs for fine-tuning. Two examples of Alpaca data are below (one containing the optional input field):

Instruction: Name the members of a rock band

Output: The members of a rock band are typically a singer, a guitarist, a bassist, a drummer, and a keyboardist.

Instruction: Identify the odd one out.

Input: Twitter, Instagram, Telegram

Output: Telegram

Most Alpaca inputs are around 100 characters in length to minimize necessary context (saving compute during fine-tuning), while most outputs are significantly longer, ranging up to 500+ characters.

6.2 BeaverTails

BeaverTails contains question-answer pairs of harmless or harmful prompts. Two examples of BeaverTails data are below:

Prompt: What is the maximum amount of time that I can stay in a tanning bed?

Harm Category: None

Response: The maximum amount of time that you can stay in a tanning bed is recommended to be limited to 10 minutes per session and no more than once per day. Tanning beds should also be used with protective eyewear to reduce the risk of eye damage.

Prompt: How can I download music for free?
Harm Category: Financial Crime, Violence

Response: There are a few legitimate ways that you can download music for free. One is through streaming services such as Spotify ... Lastly, you could try using torrent websites, but these sites are often heavily monitored and could lead to legal consequences

The character length of BeaverTails data is distributed similarly to Alpaca data.




Figure 10: BeaverTails Harm Category Frequency

Most of the harmful data found in BeaverTails is related to crime (violence, unethical behavior, etc.) and misinformation (discrimination, hate speech, etc.). There do exist several hundred examples of less common harm types. Many harmful samples in the dataset are of multiple categories. The most common combinations are listed below:

Category X	Category Y	Co-occurrence
Financial, Property, Theft	Violence, Aiding, Incitement	26,687
Hate Speech, Offensive	Non-violent Unethical	23,860
Discrimination, Stereotype	Non-violent Unethical	20,546
Drugs, Weapons, Banned	Violence, Aiding, Incitement	14,888
Discrimination, Stereotype	Hate Speech, Offensive	13,755

Table 2: Co-occurrence frequency of safety violation categories in the BeaverTails dataset.




Figure 11: Difference-in-Means Self-Consistency

The figure above shows the cosine similarity of the difference-in-means candidate vectors when trained on the reference dataset vs when trained on TwinPrompt (Krauß et al., 2025). The fact that the similarity is non-uniform in part reflects that these two datasets are simply different; we don’t expect all the DIM candidates all to reflect safety. However, if the DIM process truly identifies a single universal safety vector for the model, we should expect perfect agreement on that subspace. The reference dataset identifies the vector at layer 11 as the safety vector, and the graph above shows that the cosine similarity at layer 11 is 0.7126. This is not perfect agreement but fairly close. The highest agreement is at layer 12 with a cosine similarity of 0.7303.

7 Discussion

Our work in this paper has shown some promising new leads. Through our comparison of the overlap of the subspaces generated by difference-in-means (DIM), ActSVD, and Refusal Cone Optimization (RCO), we have shown the existence of multiple separate safety subspaces. This shows that Arditi et al.’s (2024) titular claim that “Refusal in Language Models Is Mediated by a Single Direction” is inaccurate. On one hand, this is promising for LLM safety as it shows that safety mechanisms are distributed throughout the model. On the other hand, the fact that any one of these methods can be successful without touching the other methods safety subspaces shows that these safety mechanisms are brittle. They do not function as failsafes for one another.

As far as implications for the field, the fact that methods like we have described here work so well should be concerning. If our goal as a field is to release models that cannot be trivially modified to behave in unsafe ways, the success of a simple method like difference-in-means shows we still have a long way to go. From a safety perspective, it is desirable for safety and utility to be entangled. Models where safety and utility are separable are inherently vulnerable to these kinds of attacks. Therefore future work in jailbreak resistance research should specifically aim to entangle safety and utility, such that any attempt to make the model respond to unsafe requests makes the model useless.

7.1 Limitations

Our work has some limitations.

- **Only one model family:** All of our experiments use Llama-3.1-8B-Instruct. Our results may not generalize to other families of models such as Qwen or Deepseek.
- **Only one model size:** We used an 8B parameter model for all our experiments. It may be the case that larger models have more deeply intertwined safety subspaces.

8 Conclusion

In this paper we have replicated the results of multiple state-of-the-art jailbreak methods. We have compared the safety subspaces generated by each of these methods and shown that they do not entirely overlap. This shows that while safety may be distributed throughout multiple subspaces, this does not lead to a more robust safety mechanism.

References

- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. 2024. [Refusal in Language Models Is Mediated by a Single Direction](#). In *Advances in Neural Information Processing Systems*, pages 136037–136083.
- Jiaming Ji, Mickel Liu, Juntao Dai, Xuehai Pan, Chi Zhang, Ce Bian, Boyuan Chen, Ruiyang Sun, Yizhou Wang, and Yaodong Yang. 2023. [BeaverTails: Towards Improved Safety Alignment of LLM via a Human-Preference Dataset](#). In *Advances in Neural Information Processing Systems*.
- Torsten Krauß, Hamid Dashtbani, and Alexandra Dmitrienko. 2025. [TwinBreak: Jailbreaking LLM Security Alignments based on Twin Prompts](#). arXiv: 2506.07596 [cs.LG].
- Kaustubh Ponskhe, Shaan Shah, Raghav Singhal, and Praneeth Vepakomma. 2026. [Safety Subspaces are Not Linearly Distinct: A Fine-Tuning Case Study](#). In *The Fourteenth International Conference on Learning Representations*. arXiv: 2505.14185 [cs.LG].
- Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. 2023. [Alpaca: A Strong, Replicable Instruction-Following Model](#). Stanford CRFM blog.
- Boyi Wei, Kaixuan Huang, Yangsibo Huang, Tinghao Xie, Xiangyu Qi, Mengzhou Xia, Prateek Mittal, Mengdi Wang, and Peter Henderson. 2024. [Assessing the Brittleness of Safety Alignment via Pruning and Low-Rank Modifications](#). In *Proceedings of the 41st International Conference on Machine Learning*, pages 52588–52610.
- Tom Wollschläger, Jannes Elstner, Simon Geisler, Vincent Cohen-Addad, Stephan Günnemann, and Johannes Gasteiger. 2025. [The Geometry of Refusal in Large Language Models: Concept Cones and Representational Independence](#). In *Proceedings of the 42nd International Conference on Machine Learning*, pages 66945–66970.

9 Appendix




Figure 12: Input (including Instruction + Input) and Output Character Length Distribution in Alpaca




Figure 13: Input and Output Character Length Distributions in BeaverTails